

We told the AI to think harder.


Mostly, it just wrote more.

53 test runs of ArcKit on Claude Opus 5.5 and Sonnet 5.5, about \$150 of model time. Here's what the highest effort setting actually bought.


Same requirements. Four times the wait.

One requirements document, Claude Opus 5.5, normal effort against max.


Cost



Minutes



Requirements written (average)



The real problem was the examples.

13 of the 55 ArcKit command templates we audited showed examples that contradicted their own instructions. At normal effort, the AI copies the example.

Requirements template

The instruction: every requirement needs acceptance criteria.

The examples: 1 in 30 had them.

The result: every run left them off 36 to 61 requirements, even at max.

Fix the example, not the effort.

With the templates fixed, at normal effort:

Requirements

Every requirement complete in 4 of 4 runs, on both models.

Sonnet 5.5: mentions of acceptance criteria, 52 → 132.

12 other templates

Security actions now have owners. Runbooks have rollback steps.

The 'do nothing' option now has a cost.

What we'd take from it

- 1 Check the template before raising the effort.**
Including your own customised copies.
- 2 Choose the model, not the effort level.**
In these tests it mattered more.
- 3 Measure cost and time, not just pass or fail.**
A 48-minute wait is a cost too.

Full write-up, method and every cost: arckit.org/articles